Bhavik Mehta

Seattle, USA | mehtabhavik168@gmail.com | linkedin.com/in/bhavikmehta1101 | github.com/bhavik168 | bhavikmehta.dev

Professional Summary

Software engineer who built and shipped a full-stack SaaS product using React, TypeScript, Node.js, and AWS, including OpenAI-powered recommendations and a production PostgreSQL backend. Spent two years debugging distributed backend systems at Amdocs, working hands-on with CI/CD pipelines, AWS Lambda, and service architecture. Currently finishing an MS in CS at Seattle University, looking for full-stack or backend SDE roles on teams building consumer products.

Education

Seattle University — **Sep 2024 - Jun 2026**
Masters, Computer Science — *Seattle, USA*

Savitribai Phule Pune University — **Jun 2019 - Aug 2022**
Bachelor of Engineering, Information Technology — *Pune, India*

Technical Skills

Languages: Python, TypeScript, JavaScript, Java, Go, SQL
Frontend: React, Next.js, Redux, Tailwind, TypeScript, Framer Motion
Backend: Node.js, Django, FastAPI, REST APIs, Express, GraphQL, microservices
Cloud/DevOps: AWS, Lambda, ECS, SQS, Docker, Kubernetes, CI/CD, GitHub Actions, Jenkins, Vercel
Databases: PostgreSQL, DynamoDB, MongoDB, Redis, MySQL, Kafka, Supabase, Prisma

Projects

ARIA - Autonomous Response Intelligent Assistant | github.com/ARIA — **May 2026**

- Won 1st place at AWSHacks 2026 Bedrock Track building ARIA, a real-time 911 dispatch co-pilot that surfaces one actionable recommendation in under 15 seconds via a six-agent Amazon Bedrock pipeline, giving solo dispatchers a thinking partner during live calls.
- Structured the reasoning layer so Claude Haiku 3.5 resolves location ambiguity mid-stream while Claude Sonnet 4 reconciles all specialist outputs into a single progressive recommendation card, keeping dispatchers focused on judgment instead of switching between tools.
- Built an AWS Lambda stream processor with provisioned concurrency that fires Navigation, Medical, and Hazmat agents the moment Amazon Transcribe detects domain keywords, cutting the coordination overhead that splits dispatcher attention mid-call.
- Grounded every agent output in FEMA, AHA, and MPDS protocols via Amazon Bedrock Knowledge Base with Titan Embeddings v2, with Bedrock Guardrails blocking unverified recommendations and logging every override to DynamoDB.

peekwise.ai | peekwise.ai — **Mar 2026 - Apr 2026**

- Built a competitive intelligence SaaS on Next.js 14 App Router, FastAPI, and Railway as a monorepo, using BullMQ as the sole async handoff between ingestion, AI enrichment, and orchestration layers so each cron service can scale or fail without taking the others down, cutting inter-service coupling to zero.
- Sequenced Groq free-tier LLMs for nightly batch review analysis and a frontier model for one briefing call per competitor, sizing batches and sleep intervals around free-tier TPM limits to process roughly 2,075 reviews per day at zero API cost.
- Separated hot analysis state in Redis hashes and sorted sets with 6-month TTLs from static config in Neon PostgreSQL, removing real-time AI calls from the request path and keeping every dashboard read sub-millisecond regardless of competitor count.
- Aggregated competitor review data from Google and Yelp through Apify and BrightData, then built a 3-layer deduplication and freshness-cutoff pipeline that cut redundant API calls across nightly runs and kept ingestion costs near zero at scale.

Habiwine | habiwine.com — **Mar 2025 - Jun 2025**

- Designed and shipped Habiwine solo as a full-stack SaaS with a React/TypeScript frontend, Node.js/Python REST APIs, and a PostgreSQL backend, then cut average API response latency by 25% through query optimization and middleware refactoring.
- Deployed backend services on AWS EC2 and Lambda with environment-specific configs and the frontend on Vercel, keeping the full stack stable in production without dedicated ops support.
- Wired OpenAI APIs into the backend to generate habit pattern analysis and personalized goal plans, building retry-safe error handling so AI failures never surfaced to users.
- Modeled normalized PostgreSQL and DynamoDB schemas with indexed relational paths, then tuned queries to cut data retrieval time by 30% under concurrent load.

Experience

Software Developer Intern **| Virtual Math Labs** — **Aug 2021 - Dec 2021**

- Developed modular backend components in Python and JavaScript using OOP principles and structured data modeling, cutting defect density by 20% across internal testing cycles.
- Provisioned Linux-based backend services using shell scripting, process management tools, and containerized workflows, improving environment stability across staging.
- Profiled API endpoints and refactored database queries to debug inter-service communication issues and reduce request-handling latency across critical paths.
- Contributed to peer code reviews, Git-based version control workflows, and Jenkins CI/CD pipelines, keeping builds reproducible and deployments structured.

Software Engineer **| Amdocs** — **Sep 2022 - Jul 2024**

- Automated Python and Java test frameworks to validate distributed enterprise backend services, scripting end-to-end integration tests that caught 30% more regression defects before production release.
- Diagnosed runtime and service-level failures across Linux environments by analyzing application logs, process behavior, and inter-service dependencies, speeding up root-cause identification during high-severity incidents.
- Integrated automated test suites into Jenkins CI/CD pipelines, enabling faster release cycles and more consistent deployments across distributed execution environments.
- Worked directly with backend developers and product stakeholders during defect triage, translating findings into code-level fixes that improved reliability in enterprise-grade software.